CS3017 Assessment: A Power-Law Peer-to-Peer System

Philip Hale

November 29, 2013

Contents

1 Peer Behaviour

Contents

The system consists of three types of agent: Super Peers, Ordinary Peers and a single Host Cache. The Host Cache is implemented as an agent in order to unify the connection behaviour between elements of the system. In other words, peers can connect to the Host Cache using the same mechanisms that are required to connect to each other.

It can be helpful to think of the identity of an agent as defined by the kinds of behaviours it can support. What we mean here by a behaviour is the ability to respond in some way to a particular situation or event. Some behaviours continually monitor the state of the agent in the network, whereas others are triggered by incoming messages from other agents.

The process of specifying and identifying different messages is explained in another section.

Super Peers and Ordinary Peers implement a common set of behaviours. For the sake of simplicity, the designation of peers into Super and Ordinary sets will occur during initialization of the system.

In these examples, some simplifications have been made. For example, we ignore the difference between a file and hash of the file which identifies it. When an algorithm states that it sends a file, it can be assumed that it is sending a hash of the file. In addition, this masks the fact that a peer will have a hash table of some sort in order to identify the file from the file name.

Assumptions have been made that connections will not be lost, and that messages will only be sent by peers that should send them.

In the following algorithms, the input for behaviours beginning with 'Receive' include a message name. In effect this means "the behaviour is triggered by a message with this name".

1.1 Host Cache

The Host Cache is the entry-point for the network with two responsibilities. Firstly, to maintain a list of connected peers. Secondly, to supply peers with a list of Super Peers which they can attempt to connect to

Theses behaviours can be summarised by the following piece of pseudocode. It will run when the Host Cache receives a message identified as 'Neighbours Request'.

1.1.1 Receive 'Neighbours Request'

```
input: Sender: peer
        Message Name: 'Neighbours Request'
Data: peerList: A hash of peers with value true if peer is super
1 if peer list doesn't contain peer then
2 |   add peer to peerList;
3 for connectedPeer ∈ peerList do
4 |   if peerList[connectedPeer] = true and connectedPeer ≠ peer then
5 |   |   neighbours ← neighbours ∪ peer;
6 send(peer, 'Neighbours Response', neighbours);
```

1.2 All Peers

Here we will explain the different types of behaviours implemented by both Super and Ordinary Peers. Some of these behaviours are shared between all peers, but their effects will differ depending on whether the peer is Super or Ordinary. This is a design decision which compromises simplicity of the described behaviours with simplicity of the system as a whole.[1]

1.2.1 Send 'Neighbours Request'

The first message a peer will send: asks the Host Cache for a list of peers which it can connect to in order to join the network. In addition, this message is sent whenever the knownPeers list is empty. This ensures that the peer always has some Super Peer address to connect to, in the event that they need more connections.

It only sends the request if the number of known peers is below a certain maximum. Additionally, it only sends the request if it isn't already waiting for a reply.

---

[1]The idea is analogous to the contradiction between two rules of software engineering: The Law of Demeter which aims to minimise method chaining, and the idea of class cohesion which gives an object well-bounded behaviour.

```
   Data: knownPeers: A list of neighbours known but not connected to
         the peer
   Data: hostCache: Address of the Host Cache
   Data: hasRequestedPeers: True if waiting for reply
1  if knownPeers = ∅ and ¬hasRequestedPeers then
2  |   send(hostCache, 'Neighbours Request');
3  |   hasRequestedPeer ← true;
```

1.2.2 Receive 'Neighbours Response'

A reply to the 'Neighbours Request' message. The message payload contains a list of Super Peers which the peer can try and connect to.

```
   input: Message Name: 'Neighbours Response'
          Potential Neighbours: peers
   Data: hasRequestedPeers: True if waiting for reply
   Data: knownPeers: A list of Super Peers known but not connected to
         the peer
1  knownPeers ← knownPeers ∪ peers;
2  hasRequestedPeer ← false;
```

1.2.3 Send 'Connect Request'

Asks a Super Peer to be a connection. In order to prevent blocking, the contacted Super Peer is placed in a list while we wait for a response.

The peer will keep trying to connect to more peers until the number of connected peers is greater than a given minimum. In order to limit on how many connection attempts are made, the count includes peers which we are waiting on for a reply.

```
   Data: knownPeers: A list of Super Peers known but not connected to
         the peer
   Data: connectPendingPeers: Super Peers that have been sent a
         connection request
   Data: connectedPeers: Super Peers that this peer is connected to
   Data: minPeers: Required number of connected peers.
1  if knownPeers ≠ ∅ and
   ∥connectPendingPeers ∪ connectedPeers∥ < minPeers then
2  |   send(∃p ∈ knownPeers, 'Connect Request');
```

1.2.4 Receive 'Connect Response'

A connection response is either successful or unsuccessful. An alternative approach would only send a response to the request if it is successful, but this increases the complexity of the behaviour and so has not been adopted.

A connection response from a peer removes that peer from the known peers list. If the connection is successful, the peer is also added to the list of connected peers.

```
    input: Sender: sender
           Message Name: 'Connect Response'
           Connection Successful?: connectSucceed
    Data: knownPeers: A list of Super Peers known but not connected to
          the peer
    Data: connectPendingPeers: Super Peers that have been sent a
          connection request
    Data: connectedPeers: Super Peers that this peer is connected to
  1 knownPeers ← knownPeers \ {sender};
  2 connectPendingPeers ← connectPendingPeers \ {sender};
  3 if connectSucceed = true then
  4  │  connectedPeers ← connectedPeers ∪ {sender};
```

For this demonstration, a peer will always and continually search for files they want.

1.2.5 Send 'Search Request'

```
    Data: connectedPeers: Super Peers that this peer is connected to
    Data: wantedFiles: List of files the peer wants but does not have
  1 if wantedFiles ≠ ∅ and connectedPeers ≠ ∅ then
  2  │  send(∃p ∈ connectedPeers, 'Search Request', ∃f ∈ wantedFiles);
```

1.2.6 Receive 'File Request'

Since our simulation does not send actual files, at this point we can simply return a token response to indicate that we are transferring the file. In practice, we would have to identify the file from the given identifier, and check it exists before returning etc.

```
    input: Sender: peer
           Message Name: 'File Request'
           File: file
    Data: sharedFiles: List of files the peer wants but does not have
  1 send(peer, 'File Response', 'Here is your file');
```

1.2.7 Receive 'File Response'

This completes the procedure for acquiring files in the network.

```
    input: Message Name: 'File Response'
            File: file
    Data: sharedFiles: List of files the peer wants but does not have
    Data: transferringFiles: List of files that are waiting to be sent to this
            peer.
    Data: wantedFiles: List of files the peer wants but doesn't have
  1 if file ∈ transferringFiles then
  2 │   transferringFiles ← transferringFiles \ {file};
  3 │   wantedFiles ← wantedFiles \ {file};
  4 │   sharedFiles ← sharedFiles ∪ {file};
```

1.3 Ordinary Peers

These behaviours are exclusively implemented by Ordinary Peers.

1.3.1 Send 'File List'

An Ordinary Peer is required to submit its list of shared files to its connected
Super Peer. This is so that other peers can search for files.

Since this behaviour only applies to Ordinary Peers, we can make the sim-
plifying assumption that the list of connected peers only contains a single peer.

We also ignore the possibility that a SuperPeer can become disconnected,
which is not a requirement at this stage.

One further assumption is made. We presume that the list of file hashes
that are sent to the Super Peer is small enough to fit in a single message. In
practice, this is not likely to be the case.

```
    Data: connectedPeer: Super Peer that this peer is connected to
    Data: sharedFiles: List of files the peer wants but does not have
  1 send(connectedPeer, 'File List', sharedFiled);
```

1.3.2 Receive 'Search Response'

When a peer receives a search result, the nature of the protocol means that this
must be a successful result. If there is no results found, then there is simply no
message returned. The payload of the message includes the particular file that
was found and the ID of the peer who owns it. Since in our simulation we do
not need to present results in any way, once the result is found we can contact
the peer directly to request they transfer the file.

```
    input: Message Name: 'Search Response'
            Result: file
            Peer With File: peer
    Data: transferringFiles: List of files that have been requested from
            another peer, and are waiting to be sent to this peer.
    Data: wantedFiles: List of files the peer wants but doesn't have
  1 wantedFiles ← wantedFiles \ {file};
  2 send(peer, 'File Request', file);
  3 transferringFiles ← transferringFiles ∪ {file};
```

1.4 Super Peers

These behaviours are exclusively implemented by Super Peers.

1.4.1 Receive 'Connection Request'

input: a request from a peer to join; output: connection response (y/n)

Super Peers maintain a lit of peers which are connect to them, and the files they are sharing. The first step in this interaction is for the Peer to associate itself with the Super Peer.

---

input: Sender: $peer$
      Message Name: 'Connection Request'
Data: $connectedOrdinaryPeers$: Ordinary Peers that this peer is
      connected to
Data: $ordinaryPeerLimit$: Maximum number of Ordinary Peers this
      Super Peer can connect to
1 **if** $\|connectedOrdinaryPeers\| < ordinaryPeerLimit$ **then**
2    $connectedOrdinaryPeers \leftarrow connectedOrdinaryPeers \cup \{peer\}$;
3   send($peer$, 'Connect Response', true);
4 **else**
5   send($peer$, 'Connect Response', false);

1.4.2 Receive 'File list'

The Super Peer stores the list of files for each Ordinary Peer it is connected to. This is so that it is able to handle search requests.

The files are stored as a hash that maps $fileName \rightarrow peerList$.

---

input: Sender: $peer$
      Message Name: 'File List'
      Shared Files: $fileList$
Data: $connectedOrdinaryPeers$: Ordinary Peers this peer is connected
      to
Data: $ordinaryPeersFiles$: All shared files of connected Ordinary Peers
1 **for** $file \in fileList$ **do**
2    $ordinaryPeersFiles[file]$ $peer$;

1.4.3 Receive 'Search Response'

Super Peers must behave slightly differently when receiving a Search Response message. Unlike for Ordinary Peers, sometimes this response may not be for them. This is because search results do not get delivered directly: they are passed back through intermediary peers using the Document Routing system.

It is worth noting that in our $send()$ function we pass as many arguments as we have attributes. In practice, the attributes would be wrapped into a single argument.

```
   input: Message Name: 'Search Response'
          Result: file
          Peer With File: peer
          Previous Senders of the Search Response: senderStack
1  if senderStack = ∅ then
2  |   inherit the behaviour of Ordinary Peer;
3  superPeer ← senderStack.pop();
4  send(superPeer, 'Search Response', file, peer, senderStack);
```

1.4.4 Receive 'Search Request'

When a Super Peer receives a Search Request for a file, it first checks its list of files. If the file is found, it replies with a Search Response.

If it is not found, it must pass on the message. If the $senderStack$ is empty then the sender must be an ordinary peer, so the sender is added to it. Next, it adds itself to the $senderStack$ and passes on the message to $\exists peer \in connectedOrdinaryPeers$.

The peer chosen to receive this message is picked based on the proximity of their ID when compared to the file ID. The details of this operation are not explained in this section, but it introduces a requirement of the implemented system for a consistent hashing algorithm.

```
    input: Sender: peer
           Message Name: 'Search Request'
           Requested File: file
           Previous Senders of the Search Response: senderStack
    Data: connectedOrdinaryPeers: Ordinary Peers connected to this
          Super Peer
    Data: ordinaryPeersFiles: All shared files of connected Ordinary Peers
    Data: self: This peer
1   if ordinaryPeersFiles contains file then
2   |   peerWithFile ← ∃p ∈ ordinaryPeersFiles[file];
3   |   if senderStack = ∅ then
4   |   |   send(peer, 'Search Response', file, peerWithFile)
5   |   else
6   |   |   send(senderStack.pop(), 'Search Response', file, peerWithFile,
        |   |   senderStack);
7   else
        /* We don't have the file; forward request to neighbour */
8   |   nearestNeighbour ← nearest(senderStack, file);
9   |   senderStack ← senderStack ∪ {self};
10  |   send(nearestNeighbour, 'Search Request', file, senderStack);
```

2 Protocols

Contents

The following abbreviations are used:

OPeer An Ordinary Peer
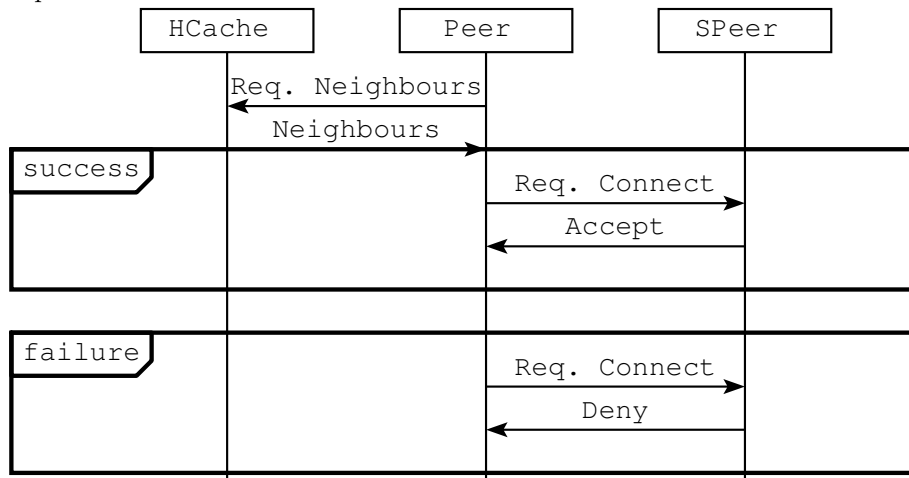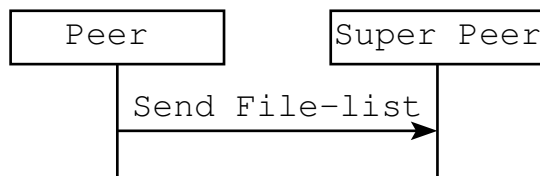
SPeer A Super Peer

Req. Send a request message

2.1 Joining the network

This involves contacting the Host Cache for the IDs of Super Peers, and requesting a connection. The connection can either be accepted or denied. If the connection is accepted, the peer sends their list of shared files to the connected Super Peer.
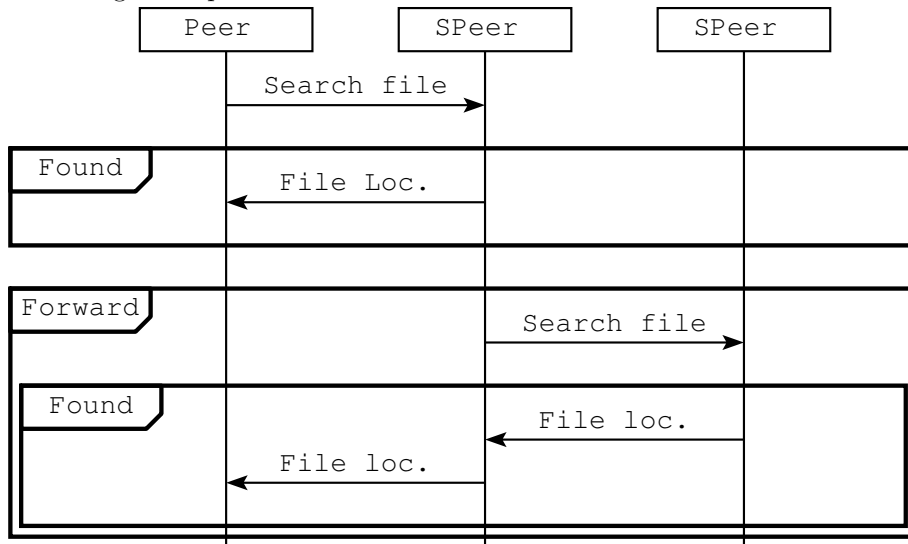

2.2 Send file-list

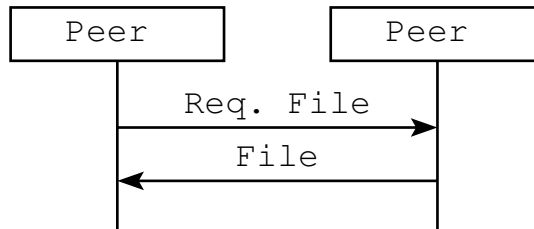Only ordinary Peers send a file list.

2.3 Searching for Files

In the following example, there are only two Super Peers in the network. The search request is forwarded throughout the network, and if at some point a Super Peer is able to locate it, the location is passed back through the network to the original requester.

```
   ┌────────┐        ┌────────┐        ┌────────┐
   │  Peer  │        │  SPeer │        │  SPeer │
   └────────┘        └────────┘        └────────┘
        │    Search file    │                │
        │──────────────────>│                │
  ┌Found┐                   │                │
  │     │     File Loc.     │                │
  │     │<──────────────────│                │
  └─────┘                   │                │

  ┌Forward┐                 │                │
  │       │          │   Search file         │
  │       │          │──────────────────────>│
  │ ┌Found┐          │                       │
  │ │     │          │      File loc.         │
  │ │     │          │<──────────────────────│
  │ │     │  File loc.     │                  │
  │ │     │<───────────────│                  │
  └─┴─────┘                │                  │
```

2.4 Acquiring Files

Once a file location has been determined, acquiring it is straightforward direct communication between two peers.

```
   ┌────────┐        ┌────────┐
   │  Peer  │        │  Peer  │
   └────────┘        └────────┘
        │   Req. File      │
        │─────────────────>│
        │      File        │
        │<─────────────────│
        │                  │
```

3 Message Definitions

- Neighbours Request
- Neighbours Response
- Search Request
- Search Response
- File Request
- File Response
- File List
- Connect Request
- Connect Response